

AMINUTEMAN TECHNOLOGIES PVT LTD

EXAMINATION TASK ASSESSMENT

Multi-Agent AI Voice System Implementation Exam

Food Delivery Platform (DoorDash-Style)

Duration: 8 Days | Production-Ready System

Problem Statement

You are tasked with building a complete multi-agent AI voice calling system for a food delivery platform similar to DoorDash. The system must handle customer orders, restaurant coordination, delivery driver management, and customer support through intelligent voice agents with real-time CRM integration and order orchestration.

System Requirements

Core Business Flows to Implement

1. Customer Order Placement (Inbound/Outbound)

- Voice-based order taking with natural conversation
- Menu navigation and item selection using voice
- Handle customizations and special dietary requests
- Address verification with Google Maps integration
- Real-time payment processing
- Order confirmation and ETA communication

2. Restaurant Order Coordination (Outbound)

- Automated order notification calls to restaurants
- Confirmation of preparation time
- Handle inventory issues (out of stock items)
- Order modification coordination
- Pickup time reminders

3. Delivery Driver Assignment (Outbound)

- Automated driver notification system
- Route optimization communication
- Pickup and delivery confirmation
- Real-time issue handling

4. Customer Support & Updates (Inbound/Outbound)

- Handle "where is my order" inquiries
- Process modification and cancellation requests
- Complaint resolution and refund processing
- Proactive delay notifications

5. Post-Delivery Follow-up (Outbound)

- Automated delivery confirmation
- Feedback and rating collection
- Issue resolution within 24 hours
- Promotional offer delivery

Technical Implementation Tasks

Section 1: Agent Architecture Design

Task 1.1: Design Multi-Agent System Create a complete agent architecture with the following agents:

1. Customer Order Agent

- Capabilities: Menu search, item addition, customization, address verification, payment
- Decision points: When to upsell, when to suggest alternatives, when to transfer
- Context requirements: Customer history, preferences, current order state
- Tool integrations: Menu search, CRM, payment gateway, maps API 2.

Restaurant Coordination Agent

- Capabilities: Order notification, time confirmation, inventory handling
- Decision points: When to suggest substitutions, when to escalate delays

- Context requirements: Restaurant profile, current queue, historical prep times
- Tool integrations: Restaurant database, order management

system 3. Driver Assignment Agent

- Capabilities: Driver search, route communication, acceptance confirmation
- Decision points: Which driver to assign, incentive calculation, reassignment logic
- Context requirements: Driver locations, ratings, availability, traffic data
- Tool integrations: GPS tracking, route optimization, driver database

4. Delivery Tracking Agent

- Capabilities: Real-time status updates, delay notification, address correction
- Decision points: When to notify customer, how to handle delivery failures
- Context requirements: Live driver location, ETA, customer availability
- Tool integrations: GPS tracking, traffic API, customer notification

5. Customer Support Agent

- Capabilities: Refund processing, complaint handling, issue escalation
- Decision points: Refund vs credit, escalate to human, compensation amount
- Context requirements: Order history, customer LTV, issue severity
- Tool integrations: Payment gateway, ticketing system, CRM

6. Post-Delivery Agent

- Capabilities: Satisfaction confirmation, feedback collection, promotion delivery
- Decision points: When to offer promotions, when to escalate issues
- Context requirements: Delivery status, customer feedback history
- Tool integrations: Review system, promotions database

Deliverable: Complete agent architecture document with interaction flows

Task 1.2: Design State Management System Implement a robust state management system that:

- Maintains conversation context across agent transitions
- Stores order state (items, pricing, delivery address)
- Preserves customer profile information
- Handles concurrent state updates safely
- Implements session persistence with appropriate TTL

Requirements:

- Use Redis for active session storage
- Design schema for session data
- Implement context handoff mechanism between agents
- Handle session recovery after disconnections
- Clean up expired sessions automatically

Deliverable: State management system design and Redis schema

Task 1.3: Agent Orchestration Logic Build the orchestration layer

- that:
- Routes incoming calls to appropriate agents
 - Handles agent-to-agent transitions seamlessly
 - Maintains conversation continuity during handoffs
 - Implements load balancing across agent instances
 - Provides failover mechanisms

Requirements:

- Intent classification system
- Agent capability registry
- Dynamic routing based on context
- Handoff summary generation
- Transition logging for analytics

Deliverable: Working orchestrator with routing logic

Section 2: Voice Infrastructure Implementation

Task 2.1: Twilio Voice Integration Implement complete Twilio voice infrastructure: **Inbound Call Handling:**

- Webhook endpoint to receive calls
- Call SID tracking and session creation
- Caller ID lookup and customer identification
- Media stream initialization

Outbound Call Handling:

- Programmatic call initiation
- Machine detection (voicemail detection)
- Call status tracking
- Retry logic for failed calls

Media Streaming:

- WebSocket connection for real-time audio
- Bidirectional audio streaming
- Call recording with consent
- Call quality monitoring

Requirements:

- Handle 100+ concurrent calls
- <500ms connection establishment time
- Proper error handling and retries
- TRAI/TCPA compliance implementation

Deliverable: Complete Twilio integration with webhooks

Task 2.2: Real-Time Audio Processing Pipeline Build audio processing system with: **Speech-to-Text (Deepgram):**

- Real-time streaming transcription
- Interim and final results handling
- Voice Activity Detection (VAD)
- Silence detection and utterance completion
- Multi-language support (English + Hindi/Tamil)

Text-to-Speech (ElevenLabs):

- Natural voice synthesis
- Streaming audio output
- Low-latency response (<500ms)
- Emotion and tone control
- Voice consistency across agents

Audio Processing Features:

- Noise reduction
- Echo cancellation
- Volume normalization
- Barge-in/interruption handling

Requirements:

- Total STT latency <300ms
- Total TTS latency <500ms
- Handle audio quality issues gracefully
- Support 8kHz and 16kHz sample rates

Deliverable: Working audio processing pipeline

Task 2.3: Interrupt and Barge-in Handling Implement sophisticated interrupt handling:

- Detect user speech while AI is speaking

- Immediately stop current TTS output
- Clear any queued responses
- Resume listening for user input
- Maintain conversation context

Requirements:

- <200ms interrupt detection time
- Graceful cancellation of ongoing speech
- No audio artifacts during interruption
- Context preservation

Deliverable: Interrupt handling system

Section 3: LLM Integration & Agent

Intelligence Task 3.1: Choose and

Implement LLM Strategy Option A:

API-Based (For 8-Day Timeline)

- Primary: Claude Sonnet 4.5 via Anthropic API
- Fallback: GPT-4o-mini for cost optimization
- Function calling for tool execution
- Streaming responses for lower latency

Option B: Self-Hosted SLM (For Future

- Scale)**
- Primary: Llama 3.3 70B Instruct
 - Lightweight: Qwen 2.5 14B for simple agents
 - Fine-tuning strategy for domain adaptation
 - vLLM for optimized inference

Requirements:

- Function calling support
- <1.5s response time

- Cost per call <\$0.05 for LLM portion
- Fallback mechanism for API failures

Decision: Document which approach you're taking and why

Deliverable: Working LLM integration with function calling

Task 3.2: Prompt Engineering for Each Agent Create production-ready system prompts for all 6 agents:

Each prompt must include:

- Clear role definition and capabilities
- Conversation style guidelines
- Tool usage instructions
- Transfer criteria to other agents
- Error handling instructions
- Example conversation flows

Quality Requirements:

- Consistent personality across agents
- Natural, conversational tone
- Efficient (minimize unnecessary back-and-forth)
- Handle edge cases and ambiguity
- Never hallucinate information

Deliverable: Complete prompt library for all agents

Task 3.3: Function/Tool Registry Implement comprehensive tool system: **Customer-facing tools:**

- get_customer_profile(phone_number)
- search_menu(restaurant_id, query, filters)
- get_item_details(item_id)
- add_to_order(item_id, quantity, customizations)

- remove_from_order(item_id)
- verify_address(address)
- calculate_total(apply_promotions)
- place_order(payment_method_id)
- get_order_status(order_id)
- process_refund(order_id, reason, amount)

Restaurant tools:

- notify_restaurant(restaurant_id, order_details)
- confirm_preparation_time(order_id, minutes)
- handle_unavailable_item(order_id, item_id)
- update_restaurant_status(restaurant_id, status)

Driver tools:

- find_available_drivers(location, radius)
- assign_driver(driver_id, order_id)
- send_pickup_details(driver_id, restaurant_info)
- update_driver_location(driver_id, coordinates)
- confirm_delivery(order_id)

Requirements:

- Type-safe function definitions
- Input validation
- Error handling
- Logging and monitoring
- Response caching where appropriate

Deliverable: Complete tool registry with implementations

Section 4: Database & Data Management

Task 4.1: Database Schema Design Design and implement PostgreSQL schema:

Required Tables:

- customers (profile, addresses, payment methods)
- restaurants (profile, menu, operating hours)
- menu_items (details, pricing, availability)
- orders (items, pricing, status, timeline)
- drivers (profile, location, availability)
- call_sessions (transcript, metrics, outcomes)
- agent_transitions (handoff tracking)

Requirements:

- Proper indexing for performance
- Foreign key relationships
- JSONB for flexible data (addresses, customizations)
- Timestamp tracking for all records
- Support for soft deletes

Deliverable: Complete database schema with migrations

Task 4.2: CRM Integration Build CRM integration

- system:
- Real-time customer data retrieval during calls
 - Order history and preferences lookup
 - Address and payment method management
 - Post-call data synchronization
 - Customer lifetime value calculation

Requirements:

- <100ms lookup time
- Caching strategy for frequent lookups
- Webhook support for real-time updates
- Error handling for CRM unavailability

Deliverable: Working CRM integration

Task 4.3: Vector Database for Semantic Search Implement Qdrant vector database for:

- Menu item semantic search ("something spicy and healthy")
- Restaurant recommendations
- FAQ and knowledge base for support agent
- Similar order suggestions

Requirements:

- Generate embeddings using OpenAI ada-002
- Create collections for menu items and FAQs
- Implement hybrid search (semantic + filters)
- Cache frequent queries

Deliverable: Vector DB integration with semantic search

Section 5: Order Management & Payment Processing

Task 5.1: Order Lifecycle Management Implement complete order flow: **Order States:**

- cart → placed → confirmed → preparing → ready → picked_up → in_transit → delivered → completed

State Transitions:

- Automated status updates
- Webhook notifications
- Real-time tracking
- Failure handling and rollbacks

Requirements:

- Atomic state transitions
- Event logging for audit trail
- Notification triggers at each state
- Handle concurrent updates

Deliverable: Order management system

Task 5.2: Payment Processing Integrate payment gateway

(Stripe/Razorpay): • Store and retrieve payment methods securely

- Process payments during order placement
- Handle payment failures and retries
- Process refunds automatically
- Support multiple payment methods

Requirements:

- PCI compliance
- Encryption for sensitive data
- Idempotency for payment requests
- Detailed transaction logging

Deliverable: Payment processing system

Task 5.3: Real-Time Order Tracking Build tracking

system: • Driver GPS location streaming

- ETA calculation with traffic data
- Proactive delay notifications
- Customer notification system (SMS/Email)
- Delivery proof capture

Requirements:

- Real-time location updates
- Accurate ETA using Google Maps API
- Geofencing for delivery confirmation
- Photo capture for proof of delivery

Deliverable: Order tracking system

Section 6: Advanced Features & Optimization

Task 6.1: Sentiment Analysis Implement real-time sentiment

- analysis:
- Analyze customer tone during calls
 - Detect frustration or dissatisfaction
 - Trigger escalation to support
 - Adjust agent response tone
 - Track sentiment metrics

Requirements:

- Real-time analysis during conversation
- <100ms processing time
- Integration with agent decision-making

Deliverable: Sentiment analysis system

Task 6.2: Call Summarization Build automatic call

- summarization:
- Generate concise summaries post-call
 - Extract key information (order details, issues)
 - Store in database for future reference
 - Use for agent training and quality assurance

Requirements:

- Accurate extraction of key details
- Structured output format
- Integration with CRM

Deliverable: Call summarization feature

Task 6.3: Cost Optimization Implement cost reduction

- strategies:
- Use cheaper TTS (Deepgram Aura) for confirmations
 - Switch to GPT-4o-mini for simple queries

- Cache frequent LLM responses
- Batch process non-urgent operations
- Optimize prompt lengths

Target: Reduce cost per call by 30%

Deliverable: Cost optimization report

Task 6.4: Multi-Language Support Implement support for 3+

- languages:
- English (primary)
 - Hindi
 - Tamil
 - Language detection
 - Seamless switching

Requirements:

- Language-specific STT/TTS models
- Translated prompts
- Bilingual menu support

Deliverable: Multi-language system

Section 7: Monitoring, Compliance & Security

Task 7.1: Monitoring Dashboard Build comprehensive

- monitoring:
- Real-time call metrics (volume, success rate, duration)
 - Agent performance tracking
 - System health monitoring
 - Cost tracking per call
 - Error rate and alerting

Tools: Prometheus + Grafana

Key Metrics:

- Calls per minute

- Average call duration
- Agent success rates
- STT/TTS/LLM latency
- Order completion rate
- Customer satisfaction score

Deliverable: Grafana dashboard

Task 7.2: Compliance

Implementation TRAI Compliance

(India):

- DND registry check before calling
- Call only between 9 AM - 9 PM IST
- Recording consent at call start
- Opt-out mechanism
- Complaint handling system

TCPA Compliance (USA):

- Prior express consent verification
- Opt-out in every call
- Respect calling hours (8 AM - 9 PM local)
- Accurate caller ID

GDPR Compliance (EU):

- PII data encryption
- Right to access data
- Right to deletion
- Consent management
- Data minimization

Deliverable: Compliance system with documentation

Task 7.3: Security Implementation Implement security

- measures:
- API rate limiting
 - Webhook signature verification
 - Data encryption at rest and in transit
 - PII anonymization in logs
 - Session token security
 - SQL injection prevention
 - XSS protection

Deliverable: Security audit checklist

Section 8: Testing & Quality Assurance

Task 8.1: Unit Testing Write comprehensive unit

- tests:
- Agent response generation
 - Tool execution
 - State management
 - Payment processing
 - Order lifecycle
 - Audio processing components

Target: >80% code coverage

Deliverable: Test suite with coverage report

Task 8.2: Integration Testing Test complete

- workflows:
- End-to-end order placement
 - Agent handoffs
 - Payment processing
 - Refund processing
 - Multi-agent coordination

- Error handling

Deliverable: Integration test suite

Task 8.3: Load Testing Perform load

testing:

- 100 concurrent calls baseline
- 500 concurrent calls stretch goal
- Identify bottlenecks
- Test failover mechanisms
- Measure response times under load

Tools: Locust or K6

Deliverable: Load test results and analysis

Task 8.4: Call Quality Testing Test conversation

quality:

- Natural conversation flow
- Accurate speech recognition
- Natural TTS output
- Interruption handling
- Context preservation
- Edge case handling

Deliverable: Call quality report

Section 9: Deployment & DevOps

Task 9.1: Containerization Dockerize all services:

- FastAPI application
- Agent orchestrator
- Audio processor
- Background workers
- Database setup

Deliverable: Docker Compose setup

Task 9.2: Kubernetes Deployment Create Kubernetes

manifests:

- Deployment configs

- Service definitions
- ConfigMaps and Secrets
- Horizontal Pod Autoscaler
- Ingress configuration
- Resource limits

Deliverable: K8s manifests

Task 9.3: CI/CD Pipeline Set up automated deployment:

- Automated testing on commit
- Docker image building
- Staging deployment
- Production deployment with approval
- Rollback capability

Tools: GitHub Actions or GitLab CI

Deliverable: Working CI/CD pipeline

Section 10: Documentation & Demo

Task 10.1: Technical Documentation Create comprehensive documentation:

- System architecture overview
- API documentation (OpenAPI/Swagger)
- Agent behavior documentation
- Database schema documentation
- Deployment guide
- Troubleshooting guide

- Runbook for operations

Deliverable: Complete documentation

Task 10.2: Demo Preparation Prepare system

demonstration: • Record 15-minute demo video

showing:

- Customer placing an order
- Restaurant coordination
- Driver assignment
- Order tracking
- Customer support interaction
- Post-delivery follow-up
- Show monitoring dashboards
- Demonstrate agent handoffs
- Show real-time state management

Deliverable: Demo video and presentation

Task 10.3: Cost Analysis Provide detailed cost

breakdown: • Per-call cost calculation

- Monthly projections at different scales (10K, 100K, 500K calls)
- Cost optimization strategies
- ROI analysis
- Break-even analysis for SLM migration

Deliverable: Cost analysis spreadsheet

Technology Stack Specification

Voice Infrastructure

- **Twilio Programmable Voice** - Call handling and media

streaming • Cost: \$0.0085/min inbound, \$0.0140/min outbound

Speech Recognition

- **Deepgram Nova-2** - Real-time STT
 - Cost: \$0.0043/min
 - Latency: <300ms

Text-to-Speech

- **ElevenLabs** (Primary) - Natural voice synthesis
 - Cost: \$0.18/1K characters
- **Deepgram Aura** (Fallback) - Cost optimization
 - Cost: \$0.015/1K characters

Language Models

Option A - API (Recommended for 8 days):

- **Claude Sonnet 4.5** - Main LLM
 - Cost: \$3/\$15 per 1M tokens
- **GPT-4o-mini** - Fallback
 - Cost: \$0.15/\$0.60 per 1M tokens

Option B - Self-Hosted (For scale):

- **Llama 3.3 70B** - Main model
- **Qwen 2.5 14B** - Lightweight
- Infrastructure: 2x A100 GPUs

Backend Framework

- **Python 3.11+ with FastAPI** - Async API server
- **Redis** - Session and cache management
- **PostgreSQL 15** - Primary database
- **Qdrant** - Vector database

Message Queue

- **RabbitMQ** - Event processing

Monitoring

- **Prometheus** - Metrics
- **Grafana** - Dashboards
- **Sentry** - Error tracking
- **ClickHouse** - Analytics

Deployment

- **Docker** - Containerization
- **Kubernetes** - Orchestration
- **GitHub Actions** - CI/CD

External APIs

- **Google Maps API** - Address verification
- **Stripe/Razorpay** - Payment processing
- **Twilio SMS** - Notifications

8-Day Implementation Schedule

Day 1: Foundation

Focus: Infrastructure setup and basic call flow

- Set up all accounts (Twilio, Deepgram, ElevenLabs, Claude)
- Initialize project structure
- Set up PostgreSQL and Redis
- Implement basic Twilio webhooks
- Test basic call answering

Deliverables:

- Working API server
- Basic call answering capability
- Database schema created

Day 2: Audio Pipeline

Focus: Real-time audio processing

- Implement Deepgram STT integration
- Build ElevenLabs TTS integration
- Create audio streaming pipeline
- Implement VAD and interruption handling

Test end-to-end audio flow

Deliverables:

- Working STT/TTS pipeline
- Interrupt handling functional
- Audio quality validated

Day 3: Agent Framework & Customer Order

Agent Focus: Core agent infrastructure

- Build base agent class
- Implement LLM service wrapper
- Create tool registry system
- Build state management
- Implement customer order agent
- Build menu search with vector DB

Deliverables:

- Base agent framework
- Customer order agent working
- Menu search functional

Day 4: Restaurant & Driver

Agents Focus: Multi-party

coordination

- Build restaurant coordination agent
- Implement driver assignment agent

Create order notification system •
Build driver search algorithm • Test
agent coordination

Deliverables:

- Restaurant agent working
- Driver agent working
- Order lifecycle automated

Day 5: Support & Tracking

Focus: Customer service

features • Build delivery
tracking agent
• Implement customer support agent •
Create refund processing
• Build proactive notification system •

Implement real-time ETA updates

Deliverables:

- Tracking agent operational
- Support agent handling issues
- Refund automation working

Day 6: Advanced Features

Focus: Polish and
optimization • Build
post-delivery agent
• Implement sentiment analysis •
Add call summarization
• Create RAG for FAQ

- Build analytics dashboard
- Optimize costs

Deliverables:

- All 6 agents complete
- Analytics working
- Cost optimizations applied

Day 7: Testing & Compliance

Focus: Quality and
compliance • Write unit tests

- Integration testing
- Load testing (100+ concurrent)
- Implement TRAI compliance •

Add security measures

- Bug fixes

Deliverables:

- Test suite complete
- Compliance features working •

System stable under load

Day 8: Documentation & Deployment

Focus: Production readiness

- Write all documentation
- Create deployment configs
- Set up CI/CD pipeline
- Deploy to staging
- Record demo video
- Prepare presentation

Deliverables:

- Complete documentation
- System deployed
- Demo video ready
- Presentation complete

Success Criteria

Your implementation will be evaluated on:

Functional Requirements

- ✓ All 6 agent types operational and coordinating seamlessly ✓ Complete end-to-end order flow working via voice ✓ Real-time CRM integration functional
- ✓ Payment processing working securely ✓ Order tracking with live updates ✓ Support and refund automation working

Performance Requirements

- ✓ Total response time <2.5 seconds ✓ System handles 100+ concurrent calls ✓ Call success rate >95% ✓ Agent handoff success rate >98% ✓ Order completion rate >80% ✓ STT accuracy >90%

Business Requirements

- ✓ Cost per call <\$0.25 ✓ Natural conversation quality ✓ Customer satisfaction score >4/5 ✓ System uptime >99%

Technical Requirements

AMINUTEMAN TECHNOLOGIES PVT LTD
EXAMINATION TASK ASSESSMENT
09-11-2025

- ✓ Clean, maintainable code ✓ Comprehensive error handling ✓ Proper logging and monitoring ✓ Security best practices implemented ✓ Scalable architecture ✓ Complete documentation

Compliance Requirements

- ✓ TRAI compliance for India ✓ Recording consent implementation ✓ PII data protection ✓ Opt-out mechanisms

Submission Requirements

Code Deliverables

1. Complete source code in Git repository
2. Database migration scripts
3. Docker and Kubernetes configs
4. CI/CD pipeline configuration
5. Environment configuration templates

Documentation Deliverables

1. System architecture document
2. API documentation (Swagger/OpenAPI)
3. Database schema documentation
4. Agent behavior documentation
5. Deployment guide
6. Troubleshooting guide
7. Operations runbook

Testing Deliverables

1. Unit test suite with coverage report
2. Integration test suite
3. Load test results and analysis
4. Call quality testing report

Demo Deliverables

1. 15-minute demo video covering all features
2. Live system demonstration capability
3. Presentation deck explaining architecture

Analysis Deliverables

1. Cost analysis spreadsheet
2. Performance benchmarks
3. Scalability analysis
4. Security audit checklist

Evaluation Criteria

Code Quality (25%)

- Clean, readable code
- Proper error handling
- Type hints and documentation
- Following best practices
- DRY principle
- Separation of concerns

Architecture (25%)

- Scalable design
- Proper component separation
- State management
- Database design
- API design
- Event-driven architecture

Functionality (25%)

- All features working
- Natural conversation flow
- Seamless agent handoffs
- Accurate speech recognition
- Quality TTS output
- Real-time responsiveness

Production Readiness (15%)

- Monitoring and alerting
- Error tracking
- Logging
- Security measures

- Compliance implementation
- Deployment automation

Documentation (10%)

- Complete technical docs
- Clear setup instructions
- API documentation
- Troubleshooting guides
- Code comments

Tips for Success

Day 1-2: Get the basics rock solid

- Don't move forward until basic call flow works perfectly
- Test audio quality extensively
- Ensure low latency from the start

Day 3-4: Focus on agent intelligence

- Invest time in prompt engineering
- Test edge cases thoroughly
- Make sure context preservation works

Day 5-6: Build for scale

- Think about performance from the start
- Implement caching strategically
- Monitor resource usage

Day 7: Test everything

- Don't skip load testing
- Test failure scenarios
- Verify compliance requirements

Day 8: Polish and present

- Good documentation is crucial

- Demo video should be professional
- Presentation should tell a story

Common Pitfalls to Avoid

1. **Latency creep** - Monitor and optimize latency from Day 1.
2. **Poor error handling** - Every external API call can fail.
3. **Weak prompts** - Invest time in prompt engineering.
4. **No caching** - Cache aggressively where appropriate.
5. **Ignoring costs** - Track costs from the beginning.
6. **Bad state management** - Context loss breaks user experience.
7. **No monitoring** - You can't fix what you can't see.
8. **Skipping tests** - Technical debt compounds quickly.
9. **Poor documentation** - Future you will thank you.
10. **Overengineering** - Focus on MVP first, iterate later.

Resources & References

Documentation

- Deepgram: <https://developers.deepgram.com/>
- ElevenLabs: <https://elevenlabs.io/docs>
- Anthropic Claude: <https://docs.anthropic.com/>
- OpenAI: <https://platform.openai.com/docs>

Tools

- FastAPI: <https://fastapi.tiangolo.com/>
- Redis: <https://redis.io/documentation>
- PostgreSQL: <https://www.postgresql.org/docs/>
- Qdrant: <https://qdrant.tech/documentation/>

Deployment

- Docker: <https://docs.docker.com/>
- Kubernetes: <https://kubernetes.io/docs/>
- GitHub Actions: <https://docs.github.com/actions>

Final Notes

This is an ambitious 8-day project that requires:

- Strong technical skills in Python and async programming
- Experience with API integrations
- Understanding of voice AI systems
- Ability to work under time pressure
- Focus on MVP while maintaining quality

Remember: The goal is not perfection, but a working, production-ready system that demonstrates all core capabilities. Prioritize functionality over features, and scalability over optimization.

Good luck with your implementation!
AMINUTEMAN TECHNOLOGIES PVT
LTD EXAMINATION TASK
ASSESSMENT 09-11-2025

Submission Deadline

Date: 18th November

Submit via:

- Git repository link
- Demo video link
- Documentation package
- Deployed system URL (if applicable)